

ENVIRONMENTAL TENDENCIES OF SALT LAKE CITY, UTAH

CORBIN APPLE, BRIDGET HYLAND, BEN STERLING

1. Introduction. It has been well established by the United States Environmental Protection Agency (EPA) and other independent organizations that average temperature is increasing across the globe[3]. This study examines monthly temperature data from two stations: Salt Lake City International Airport, Utah and Rye Patch Dam, Nevada. These stations were chosen because they are roughly at the same latitude (40.790°N and 40.498°N), longitude (111.980°W and 118.316°W), and elevation (1287.8m and 1260.3m, for Salt Lake City and Rye Patch Dam, respectively). Our data is from the National Centers for Environmental Information’s Global Summary of the Month (GSOM) dataset [2], which contains monthly data about major meteorological parameters at many locations across the country. Data for Rye Patch Dam begins in 1935, while data for Salt Lake City begins in 1948. Therefore, we only used years from 1948 to 2020 in our analysis. The dataset includes many parameters, but of particular interest to us were average temperature ($^{\circ}\text{C}$), total minutes of sunshine, number of days with thunderstorms, and total precipitation (mm) per month.

The goal of our first hypothesis is to examine whether the effects of climate change statistically differ between these two locations; if they do, this would support claims that climate change is predominantly man-made[5]. We accomplished this by calculating the monthly temperature anomaly at each location and comparing the mean temperature anomalies of the two locations. The goal of our second hypothesis is to determine predictors of temperature in Salt Lake City. We use a multiple linear regression with temperature as the response variable and minutes of sunshine, days with thunderstorms, and precipitation as the predictor variables.

2. First Hypothesis. Climate change is a well-documented phenomenon: on average, the global temperature is increasing. In this section, we test the hypothesis that the temperature increases in Salt Lake City and Rye Patch Dam are not equal.

Temperature anomaly is used to compare change in temperature over time. A mean temperature is calculated over a long period of time, and this long-term mean is subtracted from more recent observations. For example, it is known that the global average temperature was 13.9°C between 1901 and 2000. If the global average temperature in 2015 was 14.5°C , then the temperature anomaly for that year would be $14.5 - 13.9 = 0.6^{\circ}\text{C}$. When analyzing temperature data over many years, it is advisable to use a temperature anomaly rather than absolute temperature because it eliminates seasonal variation within the year, allowing for a more significant result. This is an accepted and widely used method in climate analysis [3].

The GSOM records the monthly average temperature at each location, so we determined temperature anomaly monthly. Using Excel, we calculated a January anomaly by averaging the January temperatures of each year from 1948 to 1975. Then, we subtracted this long-term mean from each January temperature from 1976 to 2020. We did the same for the other eleven months and for both locations, yielding different anomalies for each. In our data files, this is the column labeled TAVG ADJ. We used this anomaly in our comparison in place of the absolute temperature. It served to reduce the variance of each sample to produce meaningful results.

Because of the large number of observations in each sample, we invoked the central limit theorem to approximate the distribution as normal. The Shapiro-Wilk test for

normality gives a p-value of 1.66×10^{-2} for the Salt Lake City sample and 3.39×10^{-4} for the Rye Patch Dam sample, which verifies that the normal approximation is appropriate ($P < \alpha = 0.05$). The linear nature of the Q-Q plots (Figures 1 and 2) further supports our approximation.

Formally, our hypotheses are:

$$H_0 : \mu_{SLC} = \mu_{RPD} \text{ vs. } H_a : \mu_{SLC} \neq \mu_{RPD},$$

where μ_{SLC} is the mean temperature anomaly in Salt Lake City, and μ_{RPD} is the mean temperature anomaly in Rye Patch Dam. Because the two stations have similar geographical characteristics (i.e. latitude, longitude, and elevation), we determined that this is an observational matched pairs analysis and used the paired t-test at $\alpha = 0.05$ accordingly. As such, it was not necessary to determine equality of variance between the two samples. We determined the test statistic to be $t = -16.46$. $|t|$ is greater than the critical value $t_{n-1, \alpha/2} = 1.9649$ (where $n = 487$). Also, the p-value is 2.2×10^{-16} , which is less than the significance level $\alpha = 0.05$. Based on these results, we reject H_0 and conclude that the mean temperature anomaly at Salt Lake City is significantly different than the mean temperature anomaly at Rye Patch Dam. In other words, since 1975, the temperature has increased more at Salt Lake City than at Rye Patch Dam. The reason for this is not known, but previous research suggests that Salt Lake City emits a relatively large amount of carbon pollution per capita, which hastens the effects of climate change there [5]. These results, and useful statistics, are summarized in Tables 1 and 2.

We then randomly selected 10% of the data from each station to remove from the dataset to treat as missing. Although our dataset contains ample missing values on its own, we removed more values in order to evaluate the effects of missing data. Our analysis ignored months in which temperature data for Salt Lake City *or* Rye Patch Dam was missing. This resulted in a loss of total data points greater than 10%, leaving us with a new sample size (that is different for each run) for each station. Performing the Shapiro-Wilk test again, the distribution of the observations were still normal (from the p-values in Table 3 or Q-Q Plots in Figures 3 and 4). The new test statistic for a sample run was -15.03, for which $|t|$ is greater than the critical value $t_{n-1, \alpha/2} = 1.966$ (where $n = 390$). The p-value was 2.2×10^{-16} , which is less than the significance level $\alpha = 0.05$ (note that the reported p-value is the same as before because the true value is below machine precision). These statistics are summarized in Tables 3 and 4. Therefore, we reject H_0 and conclude that the mean temperature anomaly at Salt Lake City is greater than the mean temperature anomaly at Rye Patch Dam. The addition of more missing values reduced the sample size and (theoretically) resulted in a higher p-value, but did not change the outcome of the hypothesis test.

3. Second Hypothesis. The second hypothesis analyzes which environmental factors are the best predictors of temperature in Salt Lake City. According to the EPA, higher temperature results in either more or less precipitation and higher frequency of storms[6]. This study performs a multiple linear regression of temperature against minutes of sunlight, number of days with thunderstorms, and precipitation per month. In conjunction with the EPA article noted above, we also considered two key characteristics when deciding which parameters we would include in our study:

- linear relationship with temperature, and
- availability of data.

For example, we did not include total monthly snowfall or number of days with fog because they have too much missing data despite having a decent linear relationship.

The parameter with the most significant linear relationship with temperature is, not surprisingly, minutes of sunlight. A major factor we had to consider here, however, was that sunlight data is only recorded from 1965 to 2004. Thus, in order to include this parameter in our analysis, we needed to restrict our study down to this range (roughly 479 total months, or observations). Additionally, we limited the data used in our study to months for which the data was available for all four parameters listed above. Therefore, when including temperature, precipitation, and number of thunderstorms in our study, our sample size decreased from 479 months to 321 months because at least one of these parameters was missing data in the 479 months of data available between 1965 and 2004.

For each parameter identified above, we graphically examine its relation with temperature via a scatter plot with a line of best fit. As expected, there is a strong positive correlation between minutes of sunlight and temperature, as seen in Figure 5. Similarly, albeit not as strong, there also appears to be a positive correlation between number of days with thunderstorms and temperature, as seen in Figure 6. Lastly, the correlation between precipitation and temperature is negative, as seen in Figure 7, which is expected in some regions as stated by the EPA [6].

As an additional preliminary analysis, in order to obtain a strong multiple linear regression model to estimate temperature, we sought to use parameters that are not highly correlated with each other, as including additional variables that are highly correlated is more likely to simply increase model complexity as opposed to improve the model fit. Figure 8, which is a correlation matrix of the three independent parameters, suggests that the independent variables of sunlight, days with thunderstorms, and precipitation are not highly correlated, further supporting our decision to include them as parameters in our model.

At this point, we conducted the analysis on the selected parameters. Per Table 5, we can see the model has an adjusted R^2 of $0.7969 \approx 0.8$, meaning that roughly 80.0% of the variance in temperature can be explained by the three predictors. Using the coefficients from the table, we can estimate temperature as $\hat{T} = \beta_0 + X_1\beta_1 + X_2\beta_2 + X_3\beta_3$, where indices 1, 2, and 3 reference minutes of sun, days with thunderstorms, and precipitation, respectively. Substituting β values from Table 5,

$$\hat{T} \approx -5.26 + 1.04 \times 10^{-3}X_1 + 0.84X_2 - 4.80 \times 10^{-2}X_3.$$

Now that we have completed our multiple linear regression analysis on the three selected parameters, we turn to a final model comparison analysis. Here, we conduct a best subset analysis to locate the best model at each model size, including 1, 2, and 3 parameters. The best subset selection process considers all $\binom{3}{k}$ possibilities, where k is the number of parameters ($k = 1, 2, 3$). For each k , the best subset method will select the model with the lowest SSE. As can be seen in Table 6, the best models for $k = 1, 2$, and 3 are minutes of sun, minutes of sun + days with thunderstorms, and minutes of sun + days with thunderstorms + precipitation, respectively.

Now that we have our three best models at each k , we will calculate the following metrics for each model:

- adjusted R^2 ,
- AIC, and
- BIC.

The adjusted R^2 will measure what percentage of the variance in the model is explained by the parameters; this measure penalizes model complexity relative to R^2 . Akaike Information Criterion (AIC) and Bayesian Information Criterion (BIC) are both penalized-likelihood criteria, where the BIC penalizes model complexity a

bit more than AIC. Overall, the best model of the three from Table 6 will have the greatest adjusted R^2 and the lowest AIC and BIC values. Define:

- Model 1: Performs a regression against minutes of sunlight
- Model 2: Performs a regression against minutes of sunlight + days with thunderstorms
- Model 3: Performs a regression against minutes of sunlight + days with thunderstorms + precipitation

Table 7 contains these values for each model. The full model, model 3, is the best overall model to choose to estimate temperature because it has the highest adjusted R^2 and lowest AIC and BIC values.

To demonstrate our results, let us consider the following example to predict the average monthly temperature of a new observation Y^* . Notationally, the following formula enables the calculation of an $(1 - \alpha)$ -level prediction interval for Y^* .

$$Y^* \pm \hat{Y}^* + t_{n-(k+1), \alpha/2} s \sqrt{1 + \mathbf{x}^{*T} \mathbf{V} \mathbf{x}^*}$$

where \hat{Y}^* is the estimate of Y^* , $t_{n-(k+1), \alpha/2}$ is the critical value, n is the number of observations ($n = 321$), k is the number of predictors ($k = 3$), s is the estimate for σ , equivalent to \sqrt{MSE} , \mathbf{x}^* is the predictor variables for the three parameters, and $\mathbf{V} = (\mathbf{X}^T \mathbf{X})^{-1}$, where \mathbf{X} represents the matrix of the values of the predictor variables [4].

Now, given a new month with 15,000 minutes of sunlight, 4 days with thunderstorms, and 35 millimeters of total precipitation, we can calculate a 95% prediction interval. To avoid extrapolation, the parameter figures above are all within the ranges of values observed for each parameter in the 321 observations analyzed. Utilizing the prediction interval with 95% confidence, the average monthly temperature of this new month would be between 4.48°C and 19.63°C (with $\hat{Y}^* = 12.06^\circ\text{C}$).

To investigate the effects of missing data on the multiple linear regression, we removed 5% of temperature observations, 5% of sunlight observations, 5% of thunderstorm observations, and 5% of precipitation observations. We then ignored any data point for which at least one of these four observations was missing, resulting in a smaller sample size. Because of the random nature of the removal of data, the new sample size is approximately 80% of the original sample size for each run. The new adjusted R^2 is $0.797 \approx 0.8$, indicating that roughly 80% of the variation in temperature can be explained by the three predictors. This is not appreciably different from the adjusted R^2 for the full data set from 1965 to 2004, which was also approximately 80%. The new estimate for temperature is

$$\hat{T} \approx -4.79 + 1.03 \times 10^{-3} X_1 + 0.85 X_2 - 6.30 \times 10^{-2} X_3.$$

See Table 8 for the β values and p-values for all three predictor variables in the missing data case. Table 9 implies that Model 3 is once again optimal.

4. Missing Data. When data are missing at random, the value of the missing data is not related to the reason for its missingness. In this case, the complete data points are essentially a random sample of the entire dataset, so missing points can be ignored in the analysis. We evaluated the effects of data missing at random on our hypothesis test and multiple linear regression in sections 2 and 3, respectively.

Missing data is non-ignorable if the value of the missing data is related to the reason for its missingness. For example, a storm could cause data collection equipment to become inoperable, and the storm data would not be collected. In this case, it is

not possible to ignore the missing data in the analysis because the complete data is no longer a random sample, and therefore does not represent the whole of the data. The Heckman correction is a method that can correct bias from this non-random sample[1]. In the context of regression, suppose we want to fit the linear model

$$Y_i = X_i\beta + \epsilon$$

to our data. Heckman's model lets us obtain unbiased estimators of β . First, we introduce the selection equation $P(R_{yi} = 1 | X_i^s) = \Phi(X_i^s\beta^s)$, where R_{yi} is an indicator variable equal to 1 if Y_i is observed and 0 if Y_i is missing. We use maximum likelihood estimation, a method of estimating parameters by maximizing the likelihood function, to find estimates $\hat{\beta}^s$ for β^s . Using these estimates, we can compute the inverse Mills ratio $\hat{\lambda}_i = \frac{\phi(X_i^s\hat{\beta}^s)}{\Phi(X_i^s\hat{\beta}^s)}$ for each observation in the sample, where ϕ is the standard normal density and Φ is the standard normal distribution function. Finally, using the linear equation $Y_i = X_i\beta + \hat{\lambda}_i\beta_\lambda + \eta_i$, where $\eta \sim \mathcal{N}(0, \sigma_\eta^2)$, we can find estimators $\hat{\beta}$ and $\hat{\beta}_\lambda$ for β and β_λ , respectively. There also exists a one-step procedure for estimating β which generates a lower standard error, but the two-step procedure is used for faster computation.

After obtaining our estimates for β , we can draw σ_η^{2*} , β^* , and β_λ^* from the linear equation. Also, we can draw η^* from $\mathcal{N}(0, \sigma_\eta^{2*})$. Finally, for each missing Y , we can impute

$$Y_i^* = X_i\beta^* + \frac{-\phi(\widehat{X_i^s\beta^s})}{1 - \Phi(\widehat{X_i^s\beta^s})}\beta_{\lambda i}^* + \eta^*.$$

The GSOM does not record a reason for missing values, so we cannot determine whether our data are missing at random (MAR) or missing not at random (MNAR). Because we cannot determine whether our data is MAR or MNAR, this method will not apply to our specific dataset. However, if the missing data were obviously MNAR, we could use Heckman's correction to estimate the missing values in a linear regression analysis.

5. R Code Compilation Instructions. To run the R scripts, it is imperative to do so from the /data subdirectory as both reference data from a relative path. The team ran the scripts from Terminal with the "Rscript" command instead of RStudio; they could possibly be run from RStudio, but we did not test it in this manner.

If not installed already, one can install the libraries used by the scripts by entering the following commands in Terminal:

- R
- `>install.packages("ggpubr")`
- `>install.packages("ggplot2")`
- `>install.packages("latex2exp")`
- `>install.packages("car")`
- `>install.packages("corrplot")`
- `>install.packages("leaps")`


One can run the scripts by executing the following commands in Terminal:

- `cd <Base Directory>/AMS_572_Project/data`
- `Rscript Hypothesis.1.R`
- `Rscript Hypothesis.2.R`

REFERENCES

- [1] J. E. GALIMARD, S. CHEVRET, C. PROTOPODESCU, AND M. RESCHE-RIGON, *A multiple imputation approach for MNAR mechanisms compatible with Heckman's model*, *Statistics in Medicine*, 35(2016), pp. 2907-2920.
- [2] J. H. LAWRIK, R. RAY, S. APPLEQUIST, B. KORZENIEWSKI, AND M. J. MENNE, *Global summary of the month (gsom) dataset, version 1*, 2016, <https://doi.org/10.7289/V5QV3JJ5> (accessed 2020-11-14).
- [3] NASA, *The exploring the environment project*, <http://ete.cet.edu/gcc/?/globaltemp-anomalies/>.
- [4] A. C. TAMHANE AND D. D. DUNLOP, *Statistics and Data Analysis from Elementary to Intermediate*, Prentice-Hall, Inc., Upper Saddle River, NJ, USA, 2000.
- [5] THE WHITE HOUSE, *The threat of carbon pollution: Utah*, <https://obamawhitehouse.archives.gov/sites/default/files/docs/state-reports/climate/Utah%20Fact%20Sheet.pdf>.
- [6] UNITED STATES ENVIRONMENTAL PROTECTION AGENCY, *What climate change means for utah*, <https://19january2017snapshot.epa.gov/sites/production/files/2016-09/documents/climate-change-ut.pdf>.

6. Appendix: Tables and Figures.


FIG. 1. *Q-Q Plot in Salt Lake City*

	$\Delta Temp$ Salt Lake City	$\Delta Temp$ Rye Patch
μ	1.1229	0.1317
σ	1.9765	1.9226
n	487	487
$P_{Shapiro}$	0.0166	0.0003

TABLE 1
Temperature Difference Statistics

Parameter	Value
t	-16.46
df	486
P	2.2×10^{-16}
Conf. Interval	$(-1.1095, -0.8728)$

TABLE 2
Paired t Test Results for Temperature Difference

FIG. 2. *Q-Q Plot in Rye Patch*

	$\Delta Temp$ Salt Lake City	$\Delta Temp$ Rye Patch
μ	1.1282	0.1457
σ	1.9850	1.9051
n	390	390
$P_{Shapiro}$	0.0238	0.0005



TABLE 3
Temperature Difference Statistics for Missing Values

Parameter	Value
t	-15.03
df	389
P	2.2×10^{-16}
Conf. Interval	$(-1.1109, -0.8539)$

TABLE 4
Paired t Test Results for Temperature Difference for Missing Values


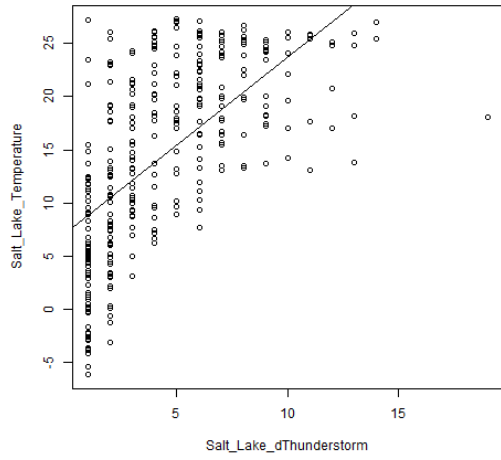
	Temperature	Intercept	Minutes of Sun	Days with Thunderstorms	Precipitation
$P_{Shapiro}$	2.662×10^{-8}	—	6.845×10^{-9}	9.438×10^{-15}	2.817×10^{-12}
β	—	-5.261	1.041×10^{-3}	0.8471	-4.800×10^{-2}
$P(> t)$	—	2.30×10^{-8}	$< 2.00 \times 10^{-16}$	$< 2.00 \times 10^{-16}$	$< 2.91 \times 10^{-7}$

TABLE 5
Multiple Linear Regression for Predicting Temperature (Adjusted $R^2 = 0.7969$, $n = 321$)

FIG. 3. *Q-Q Plot in Salt Lake City with 10% missing data*FIG. 4. *Q-Q Plot in Rye Patch with 10% missing data*

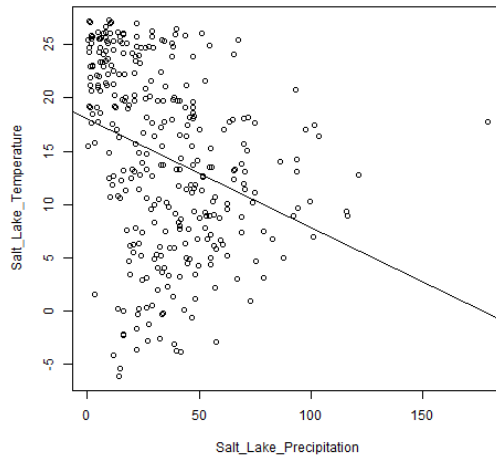

# Variables	Minutes of Sun	Days with Thunderstorms	Precipitation
1	*		
2	*	*	
3	*	*	*

TABLE 6
Optimal Subset selection for Multiple Linear Regression

FIG. 5. *Salt Lake City Temperature versus Sun*FIG. 6. *Salt Lake City Temperature versus Days with Thunderstorms*

	Model 1	Model 2	Model 3
AIC	1875.498	1805.823	1781.139
BIC	1886.812	1820.909	1799.996
Adjusted R^2	0.7258	0.78	0.7969

TABLE 7
Model Comparisons for Multiple Linear Regression

FIG. 7. *Salt Lake City Temperature versus Precipitation*FIG. 8. *Correlation plot of chosen dependent variables*

	Temperature	Intercept	Minutes of Sun	Days with Thunderstorms	Precipitation
$P_{Shapiro}$	3.259×10^{-7}	—	2.592×10^{-7}	2.011×10^{-13}	1.88×10^{-8}
β	—	-4.795	1.037×10^{-3}	0.8541	-6.308×10^{-2}
$P(> t)$	—	4.85×10^{-6}	$< 2.00 \times 10^{-16}$	$< 2.00 \times 10^{-16}$	$< 1.94 \times 10^{-8}$

TABLE 8

Multiple Linear Regression for Predicting Temperature with Missing Data (Adjusted $R^2 = 0.797$, $n = 260$)

	Model 1	Model 2	Model 3
AIC	1530.985	1481.879	1451.769
BIC	1541.667	1496.121	1469.572
Adjusted R^2	0.7226	0.7713	0.797

TABLE 9

Model Comparisons for Multiple Linear Regression with Missing Data